

---

**INDIAN INSTITUTE OF INFORMATION  
TECHNOLOGY RAICHUR**  
(IIITR)

---

# **Physics-Informed Neural Networks and Deep Ritz Method for Biharmonic Equation Solution**

*A Comparative Study of Two Numerical Examples*

## **Team Members**

**Harsh Tripathi**  
AD23B1021

**Rudra Pratap Singh**  
AD23B1047

**Lakshay Mittal**  
AD23B1027

**Sai Pranay Omkaram**  
AD23B1048

**Atla Smaran Reddy**  
CS23B1011

**L. Nitish Raman**  
CS23B1041

Under the Guidance of

**Dr. Ramesh Chandra Sau**

**December 2025**

# Physics-Informed Neural Networks and Deep Ritz Method for Biharmonic Equation Solution

## A Comparative Study of Two Numerical Examples

---

### Executive Summary

This report presents a comprehensive analysis of two computational approaches for solving fourth-order biharmonic equations: Physics-Informed Neural Networks (PINN) and the Deep Ritz Method (DRM). We investigate two distinct test cases with different boundary conditions and exact solutions, comparing accuracy, convergence behavior, and computational efficiency across methodologies.

### Key Findings:

- Both PINN and DRM successfully solve fourth-order PDEs with simply supported boundary conditions
  - Two-stage optimization (Adam → L-BFGS) consistently outperforms single-optimizer approaches
  - Example 1 (trigonometric solution) achieves relative L2 errors on the order of  $10^{-3}$
  - Example 2 (polynomial solution) demonstrates superior accuracy with errors approaching  $10^{-5}$
  - DRM formulation requires careful boundary condition weighting for optimal performance
- 

## 1. Problem Formulation

### 1.1 Biharmonic Equation

We consider the fourth-order biharmonic equation on the unit square  $\Omega = (0,1)^2$ :

$$\begin{aligned}\Delta^2 u &= f \quad \text{in } \Omega \\ u &= g_1 \quad \text{on } \partial\Omega \\ \Delta u &= g_2 \quad \text{on } \partial\Omega\end{aligned}$$

where:

- $\Delta^2 = \Delta(\Delta u)$  is the biharmonic operator  $f$  is the source term  $g_1$  and
- $g_2$  are Dirichlet boundary conditions on  $u$  and its Laplacian
- 

### 1.2 Test Cases

#### Example 3.1: Trigonometric Solution

#### Exact Solution:

$$u(x,y) = (1/(2\pi^2)) \sin(\pi x) \sin(\pi y)$$

### Source Term:

$$f(x,y) = \pi^2 \sin(\pi x) \sin(\pi y)$$

### Boundary Conditions:

- P1:  $u = g_1, g_2 = \partial u / \partial n$  on  $\partial \Omega$
- P2:  $u = g_1, g_2 = \partial^2 u / \partial n^2$  on  $\partial \Omega$
- P3:  $u = g_1, g_2 = \Delta u$  on  $\partial \Omega$  (Simply Supported)
- P4:  $u = \partial u / \partial n, g_2 = \partial \Delta u / \partial n$  on  $\partial \Omega$

**Implementation:** Problem P3 with simply supported boundaries

### Example 3.2: Polynomial Solution

#### Exact Solution:

$$u(x,y) = x^2 y^2 (1-x)^2 (1-y)^2$$

**Source Term:** Computed from  $\Delta^2 u$  where  $u = A(x) \cdot B(y)$  with  $A(t) = B(t) = t^2(1-t)^2$ :

$$f(x,y) = 24[A(x) + B(y)] + 2 \cdot A''(x) \cdot B''(y)$$

**Boundary Conditions:** Problem P3 (simply supported)

## 2. Methodology

### 2.1 Physics-Informed Neural Networks (PINN)

#### Architecture:

- Input:  $(x, y) \in [0,1]^2$
- Hidden layers: 4 layers  $\times$  128 neurons
- Activation: Tanh
- Output:  $u(x,y)$
-

## Weight initialization: Xavier uniform **Loss**

### Function:

$$L_{PINN} = \lambda_{int} \cdot L_{int} + \lambda_{bc} \cdot L_{bc}$$

where:

- **Interior Loss:**  $L_{int} = \text{MSE}(\Delta^2 u_{\theta} - f)$
- **Boundary Loss:**  $L_{bc} = \text{MSE}(u_{\theta} - g_1) + \text{MSE}(\Delta u_{\theta} - g_2)$

### Training Data:

- Interior points: 10,000-20,000 randomly sampled
- Boundary points: 8,000-10,000 uniformly distributed

### Optimization Strategy:

1. **Phase 1 - Adam:** 5,000 epochs, learning rate  $10^{-3}$
2. **Phase 2 - L-BFGS:** max\_iter=2,000, strong Wolfe line search

## 2.2 Deep Ritz Method (DRM)

### Architecture:

- Input:  $(x, y) \in [0,1]^2$
- Hidden layers: Flexible (3-5 layers, 64-128 neurons per layer)
- Activation: tanh for energy formulation OR Tanh for strong form matching
- Output:  $u(x,y)$
- Weight initialization: Xavier normal/uniform

### Energy Functional:

$$J[u] = \frac{1}{2} \int_{\Omega} |\Delta u|^2 d\Omega - \int_{\Omega} fu d\Omega$$

### Empirical Loss (P3 Formulation):

$$L_{\lambda}(\theta) = (|\Omega|/N) \sum_i B_i u(X_i) + (|\partial\Omega|/M) \sum_j [-g_2(Y_j) \partial u / \partial n(Y_j) + \lambda/2 (u - g_1)^2(Y_j)]$$

where:

- $B_{\mathbf{u}}(\mathbf{X}) = \frac{1}{2}|\Delta \mathbf{u}(\mathbf{X})|^2 - \mathbf{f}(\mathbf{X})\mathbf{u}(\mathbf{X})$  is the bilinear form
- $\partial \mathbf{u} / \partial \mathbf{n}$  is the normal derivative on boundary  $\lambda$  is the
- penalty parameter for boundary enforcement
- $\mathbf{X}_i$  are interior sample points
- $\mathbf{Y}_j$  are boundary sample points

### Advanced Sampling Strategy:

- **Stratified interior sampling:** Uniform random distribution
- **Biased boundary sampling:** Higher density at edge centers using beta-like transformation

$$t_{\text{biased}} = t(1-\beta) + 0.5\beta + (t-0.5)(1-\beta)\text{sign}(t-0.5)$$

where  $\beta = 0.3$  controls concentration at edge centers

- Motivation: Errors typically concentrate at boundaries and corners

### Normal Derivative Computation:

Exact geometric normals computed per edge:

- Left edge ( $x=0$ ):  $\mathbf{n} = (-1, 0)$
- Right edge ( $x=1$ ):  $\mathbf{n} = (1, 0)$
- Bottom edge ( $y=0$ ):  $\mathbf{n} = (0, -1)$
- Top edge ( $y=1$ ):  $\mathbf{n} = (0, 1)$

1) Then:  $\partial \mathbf{u} / \partial \mathbf{n} = \nabla \mathbf{u} \cdot \mathbf{n}$

### Training Data:

- Example 1: Interior = 8,000, Boundary = 3,000
- Example 2: Interior = 8,000, Boundary = 3,000
- Test grid:  $51 \times 51$  to  $100 \times 100$

### Optimization Strategy:

#### Single-Stage (Examples 1 & 2):

- Adam only: 3,000 epochs
- Learning rate:  $10^{-4}$  (Example 1),  $10^{-2}$  (Example 2)
- Gradient clipping: `max_norm = 1.0`
- Loss clamping:  $[-10^8, 10^8]$  to prevent NaN (bound them)
- Learning rate scheduler: StepLR (Example 2 only,  $\gamma=0.9$  every 200 epochs)

### **Two-Stage (Alternative):**

1. **Phase 1 - Adam:** 5,000-10,000 epochs, learning rate  $10^{-3}$

2. **Phase 2 - L-BFGS:** `max_iter=200, max_eval=250`

### **Stability Enhancements:**

- Gradient clipping prevents exploding gradients in deep networks
  - Loss clamping handles numerical overflow in Laplacian computation
  - Conservative learning rate balances convergence speed and stability
- 

## **3. Numerical Results**

### **3.1 Example 1: Trigonometric Solution**

#### **PINN Results**

#### **Training Configuration:**

- Network:  $[2 \rightarrow 128 \rightarrow 128 \rightarrow 128 \rightarrow 128 \rightarrow 1]$
- Interior points: 10,000
- Boundary points: 8,000
- Loss weights:  $\lambda_{\text{int}} = 1.0, \lambda_{\text{bc}} = 2.0$

#### **Convergence:**

- Final Adam loss:  $\sim 10^{-3}$
- Final L-BFGS loss:  $\sim 10^{-4}$
- Total training time:  $\sim 200\text{-}300$  seconds

#### **Error Metrics:**

Metric	Value	Description
Relative L2	$7 \times 10^{-2}$	Function approximation error
Relative H1	$\times 10^{-1}$	Including first derivatives
Relative H2	0.39	Including second derivatives

## DRM Results

### Training Configuration:

- Network:  $[2 \rightarrow 64 \rightarrow 64 \rightarrow 64 \rightarrow 64 \rightarrow 64 \rightarrow 1]$  (5 hidden layers)
- Activation: Tanh (matching PINN architecture)
- Interior points: 8,000 (stratified sampling)
- Boundary points: 3,000 (biased sampling toward edge centers)
- Penalty parameter:  $\lambda = 1,000$
- Learning rate:  $10^{-4}$  (conservative for stability)
- Gradient clipping:  $\text{max\_norm} = 1.0$

### Advanced Features:

- **Stratified boundary sampling:** Higher density at edge centers where error typically concentrates
- **Normal derivative computation:** Exact geometric normals for each boundary edge
- **Loss clamping:** Prevents NaN propagation (range:  $[-10^8, 10^8]$ )
- **Gradient stabilization:** Clipping prevents exploding gradients in deep network

### Convergence:

- Training: 3,000 epochs ( $\sim 120\text{-}180$  seconds on GPU)
- Final loss:  $\sim 10^{-5}$  to  $10^{-6}$
- Stable convergence with smooth loss decay

### Error Metrics:

Metric	Value	Performance
L2 error	$\sim 10^{-3}$	Excellent
Relative L2	$\sim 3 \times 10^{-2}$	High accuracy
Maximum error	$\sim 3.4 \times 10^{-3}$	Uniform accuracy
Mean absolute error	$\sim 7 \times 10^{-4}$	Very good

### Key Improvements Over Initial Implementation:

- 10× improvement in L2 error through deeper architecture
- Better boundary condition satisfaction via biased sampling
- More stable training with gradient clipping and loss clamping
- Matches PINN architecture for fair comparison

## 3.2 Example 2: Polynomial Solution

### PINN Results

#### Training Configuration:

- Network:  $[2 \rightarrow 128 \rightarrow 128 \rightarrow 128 \rightarrow 128 \rightarrow 1]$
- Interior points: 20,000
- Boundary points: 10,000
- Loss weights:  $\lambda_{\text{int}} = 1.0, \lambda_{\text{bc}} = 1.0$
- Adam learning rate:  $10^{-2}$  (increased from Example 1)

#### Convergence:

- Best Adam loss:  $\sim 10^{-4}$  to  $10^{-5}$
- Final L-BFGS loss:  $\sim 10^{-6}$  to  $10^{-7}$
- Total training time: ~250-350 seconds

#### Error Metrics:

Metric	Value	Improvement over Example 1
Relative L2	$1-3 \times 10^{-4}$	10× better
Relative H1	$2-5 \times 10^{-4}$	15× better
Relative H2	$1-3 \times 10^{-3}$	5-10× better

## DRM Results

### Training Configuration:

- Network:  $[2 \rightarrow 64 \rightarrow 64 \rightarrow 64 \rightarrow 64 \rightarrow 64 \rightarrow 1]$  (5 hidden layers)
- Activation: Tanh
- Interior points: 8,000
- Boundary points: 3,000
- Penalty parameter:  $\lambda = 10,000$  (higher than Example 1)
- Learning rate:  $10^{-2}$  (aggressive for polynomial smoothness)
- Learning rate scheduler: StepLR ( $\gamma=0.9$  every 200 epochs)

### Advanced Features:

- **Uniform boundary sampling:** Polynomial BCs satisfied naturally at boundaries
- **Higher penalty parameter:**  $\lambda = 10,000$  vs 1,000 for stronger BC enforcement
- **Aggressive learning rate:**  $10^{-2}$  with scheduling for faster convergence
- **Weighted boundary loss:**  $0.01 \times$  coefficient balances domain and boundary
- **Analytical bilaplacian:** Source term computed exactly from derivatives

**Source Term Computation:** For  $u = x_1^2 x_2^2 (1-x_1)^2 (1-x_2)^2$ , the bilaplacian is computed analytically:

$$\begin{aligned}\Delta^2 u = & 12x_2^2(1-x_2)^2(6x_1^2-6x_1+1) + 24x_1(1-2x_1)(6x_2^2-6x_2+1) + \\ & 12x_1^2(1-x_1)^2(6x_2^2-6x_2+1) + 12x_1^2(1-x_1)^2(6x_1^2-6x_1+1)\end{aligned}$$

### Convergence:

- Training: 3,000 epochs (~180-240 seconds on GPU)
- Learning rate decay:  $10^{-2} \rightarrow \sim 10^{-4}$  over training
- Final loss:  $\sim 10^{-6}$  to  $10^{-7}$
- Smooth monotonic convergence

### Error Metrics:

Metric	Value
L2	~2.8 *
Relative	~1.8
Maximum	~1.7
Mean absolute	~1.12

### Key Differences from Example 1:

- **No biased sampling needed:** Polynomial BCs naturally zero at boundaries
  - **Higher penalty (10 $\times$ ):** Stronger enforcement for polynomial solution
  - **Aggressive LR with scheduling:** Exploits polynomial smoothness
  - **Weighted boundary loss:** 1 $\times$  coefficient prevents boundary over-fitting
  - **Better accuracy:** ~10% improvement over Example 1 due to solution smoothness
- 

## 4. Comparative

**4.1 Key Insight:** DRM achieves superior accuracy through:

1. **Deeper architecture** (5 layers vs 4) for better representation
2. **Smarter sampling** with boundary bias where errors concentrate
3. **Better stability** through gradient clipping and loss clamping
4. **Proper normal derivative** computation using geometric normals

**Training Strategy Differences:**

Aspect	Example 1 (Trig)	Example 2 (Poly)
Sampling	Biased ( $\beta=0.3$ )	Uniform
Learning Rate	$10^{-4}$ conservative	$10^{-2}$ aggressive
LR Scheduling	None	StepLR ( $\gamma=0.9/200$ )
Penalty $\lambda$	1,000	10,000
Boundary Weight	1.0	0.01
Epochs	3,000	3,000
Final L2	$\sim 10^{-5}$ to $10^{-6}$	$\sim 10^{-6}$ to $10^{-7}$

## 5. Key Observations

### 5.1 Architecture Design

**PINN:**

- Deep networks (4 hidden layers) handle fourth-order derivatives well
- Tanh activation provides smooth, differentiable outputs
- Xavier initialization prevents vanishing/exploding gradients

**DRM:**

- **Deeper architecture** (5 hidden layers) improves representation capacity
- **Tanh activation** provides consistency with PINN comparison
- Alternative  $\sin(\cdot)$  activation captures oscillatory behavior naturally
- **Gradient clipping** essential for stability in deep networks with fourth-order derivatives
- **Loss clamping** prevents NaN propagation in Laplacian computations

### 5.2 Boundary Condition Treatment

**Critical Finding:** BC enforcement strategy significantly impacts accuracy

**PINN Approach:**

- Soft penalty through loss weighting ( $\lambda_{bc} = 2.0$ )
- Simultaneous enforcement of  $u = g_1$  and  $\Delta u = g_2$
- Uniform boundary sampling

## DRM Approaches:

### Original Implementation:

- Energy-based with separate BC weights ( $w_{bc1} = 100$ ,  $w_{bc2} = 65-75$ )
- Three separate loss terms: energy, BC1, BC2
- Uniform boundary sampling

### Current Implementation:

- **Integrated penalty method:**  $\lambda$  parameter directly in energy functional
- **Normal derivative form:**  $\partial u / \partial n$  computed using geometric normals
- **Biased boundary sampling:** Higher density at edge centers
  - Beta-like transformation
  - Bias parameter  $\beta = 0.3$  concentrates samples near  $t = 0.5$
  - Rationale: Errors peak at boundaries and corners
- **Penalty parameter:**  $\lambda = 1e3$  for strong BC enforcement

### Comparison:

Method	Boundary Loss	Sampling
PINN	MSE-based soft	Uniform
DRM Original	Weighted MSE	Uniform
DRM Current	Penalty + normal deriv	Biased

## 5.3 Loss Function Behavior

### DRM:

- **Smoothen convergence** due to gradient clipping
- **More stable** with loss clamping preventing NaN
- **Domain and boundary losses** separately tracked
- Energy term  $B_u(X) = \frac{1}{2}|\Delta u|^2 - fu$  naturally balances competing objectives
- Typical convergence pattern:
  - Epochs 0-500: Rapid descent
  - Epochs 500-1500: Steady improvement
  - Epochs 1500-3000: Fine-tuning

### **Loss Component Analysis:**

$$\begin{aligned}\text{Total Loss} &= \text{Domain Loss} + \text{Boundary Loss} \\ &= (|\Omega|/N) \sum B_u(X_i) + (|\partial\Omega|/M) \sum [\text{boundary terms}]\end{aligned}$$

### **Clamping Strategy:**

- **Purpose:** Prevent NaN from overflow in  $|\Delta u|^2$  for large gradients
  - **Range:**  $[-10^8, 10^8]$  for all loss components
  - **Impact:** Enables training with aggressive learning in early epochs
  - **Trade-off:** May clip legitimate large values, but improves stability
- 

## **6. Visualizations**

### **Visualization Insights:**

1. **Biased sampling effectiveness:** Compare Example 1 original vs current - current shows uniform error while original has boundary concentration
  2. **Architecture impact:** Deeper networks (5 layers) reduce interior errors significantly
  3. **Convergence stability:** Current DRM shows smooth monotonic decay vs. oscillatory behavior in original
  4. **Solution type matters:** Example 2 shows uniformly low error regardless of sampling strategy
  5. **Learning rate scheduling:** Example 2 benefits from aggressive LR with decay, visible in smooth loss curves
-

## 7. Hyperparameter Sensitivity

### 7.1 Critical Parameters

PINN:

Parameter	Range Tested	Optimal	Sensitivity
$\lambda_{\text{int}}$	[0.5, 2.0]	1.0	Low
$\lambda_{\text{bc}}$	[1.0, 5.0]	2.0	Medium
Adam LR	$[10^{-4}, 10^{-2}]$	$10^{-3}$	Medium
Network width	[64, 256]	128	Low
Network depth	[3, 5]	4	Medium

DRM (Current - Example 1):

Parameter	Range Tested	Optimal	Sensitivity
$\lambda$ (penalty)	[100, 10,000]	1,000	High
Adam LR	$[10^{-5}, 10^{-3}]$	$10^{-4}$	High
Boundary bias $\beta$	[0.1, 0.5]	0.3	Medium
Grad clip norm	[0.5, 2.0]	1.0	Medium
Loss clamp	$[10^6, 10^{10}]$	$10^8$	Low
Interior samples	[5k, 15k]	8,000	Low
Boundary samples	[1k, 5k]	3,000	Medium
Network depth	[3, 6]	5	High
Epochs	[1k, 5k]	3,000	Medium

DRM (Current - Example 2):

Parameter	Range Tested	Optimal	Sensitivity
$\lambda$ (penalty)	[1k, 50k]	10,000	High
Adam LR	$[10^{-3}, 10^{-1}]$	$10^{-2}$	Medium
LR decay $\gamma$	[0.7, 0.95]	0.9	Low
LR decay step	[100, 500]	200	Low
Boundary weight	[0.001, 0.1]	0.01	High
Grad clip norm	[0.5, 2.0]	1.0	Medium
Interior samples	[5k, 15k]	8,000	Low
Boundary samples	[1k, 5k]	3,000	Low
Network depth	[3, 6]	5	Medium
Epochs	[1k, 5k]	3,000	Low

### DRM (Original):

Parameter	Range Tested	Optimal	Sensitivity
w_bc1	[50, 200]	100	High
w_bc2	[40, 100]	65-75	High
Adam LR	$[10^{-4}, 10^{-2}]$	$10^{-3}$	Medium
Adam epochs	[5k, 20k]	10k	Medium

### Key Observations:

#### 1. Learning Rate: Problem-dependent optimization

- **Example 1:**  $10^{-4}$  conservative rate prevents oscillations in trigonometric solution
- **Example 2:**  $10^{-2}$  aggressive rate with scheduling exploits polynomial smoothness
- Deeper networks more sensitive to step size
- LR scheduling critical for Example 2 convergence

#### 2. Boundary Bias ( $\beta$ ): Novel parameter for oscillatory solutions **Example 1:** $\beta = 0.3$ critical for accuracy

- (10-100× improvement)
- **Example 2:**  $\beta = 0$  (uniform sampling) sufficient, polynomial BCs naturally satisfied
- $\beta > 0.5$ : Over-concentration, interior under-represented

#### 3. Boundary Loss Weighting: Critical for polynomial solutions

- **Example 1:** Weight = 1.0 (equal balance)
- **Example 2:** Weight = 0.01 prevents boundary over-fitting
- Too high weight in Example 2 causes boundary over-emphasis, degrades interior

#### 4. Network Depth: Critical for accuracy

- 3 layers: Adequate for energy minimization (original)
- 4 layers: Good for strong form enforcement (PINN)
- 5 layers: Optimal for deep Ritz with complex solutions (current)
- 6+ layers: Diminishing returns, training instability

#### 5. Learning Rate Scheduling:

- **Not needed for Example 1:** Conservative LR sufficient
- **Essential for Example 2:** StepLR enables aggressive initial learning with refinement
- $\gamma = 0.9$  every 200 epochs provides smooth decay

## 8. Limitations and Challenges

### 8.1 Boundary Condition Enforcement

#### Issues:

- Soft constraints may violate BCs slightly
- Trade-off between interior accuracy and BC satisfaction
- Second BC ( $\Delta u = g_2$ ) harder to enforce than first

#### Solutions:

- Higher boundary point density
- Adaptive weight scheduling
- Hard constraint parameterization (future work)

### Contribution Statement

**Project Title:** Physics-Informed Neural Networks and Deep Ritz Method for Biharmonic Equation Solution  
**Team Members:**

1. **Lakshay Mittal (AD23B1027)**
2. **Rudra Pratap Singh (AD23B1047)**

- 3. Harsh Tripathi (AD23B1021)**
- 4. Sai Pranay Omkaram (AD23B1048)**
- 5. Atla Smaran Reddy (CS23B1011)**
- 6. L. Nitish Raman (CS23B1041)**

#### **Equal Contribution Breakdown:**

All team members contributed equally to the following phases of the project:

- **Methodology Implementation:** Collaboratively implemented both PINN and DRM frameworks, including architecture design, loss function formulation, and optimization strategies.
- **Numerical Experimentation:** Equally involved in running simulations, tuning hyperparameters, and analyzing results for both trigonometric and polynomial examples.
- **Code Development & Debugging:** Shared responsibility in writing, testing, and refining the codebase, ensuring reproducibility and accuracy.
- **Result Analysis & Visualization:** Jointly interpreted numerical results, generated plots, and derived insights from error metrics and convergence behavior.
- **Report Writing & Documentation:** Co-authored all sections of the final report, including abstract, methodology, results, discussion, and conclusions.